

DATA CUBES

Hari Sundaram

hs1@illinois.edu


<http://sundaram.cs.illinois.edu>

adapted from slides by Jiawei Han and Kevin Chang



BASIC CONCEPTS

Methods Advanced Processing Multidimensional Summary



Important!

academic integrity

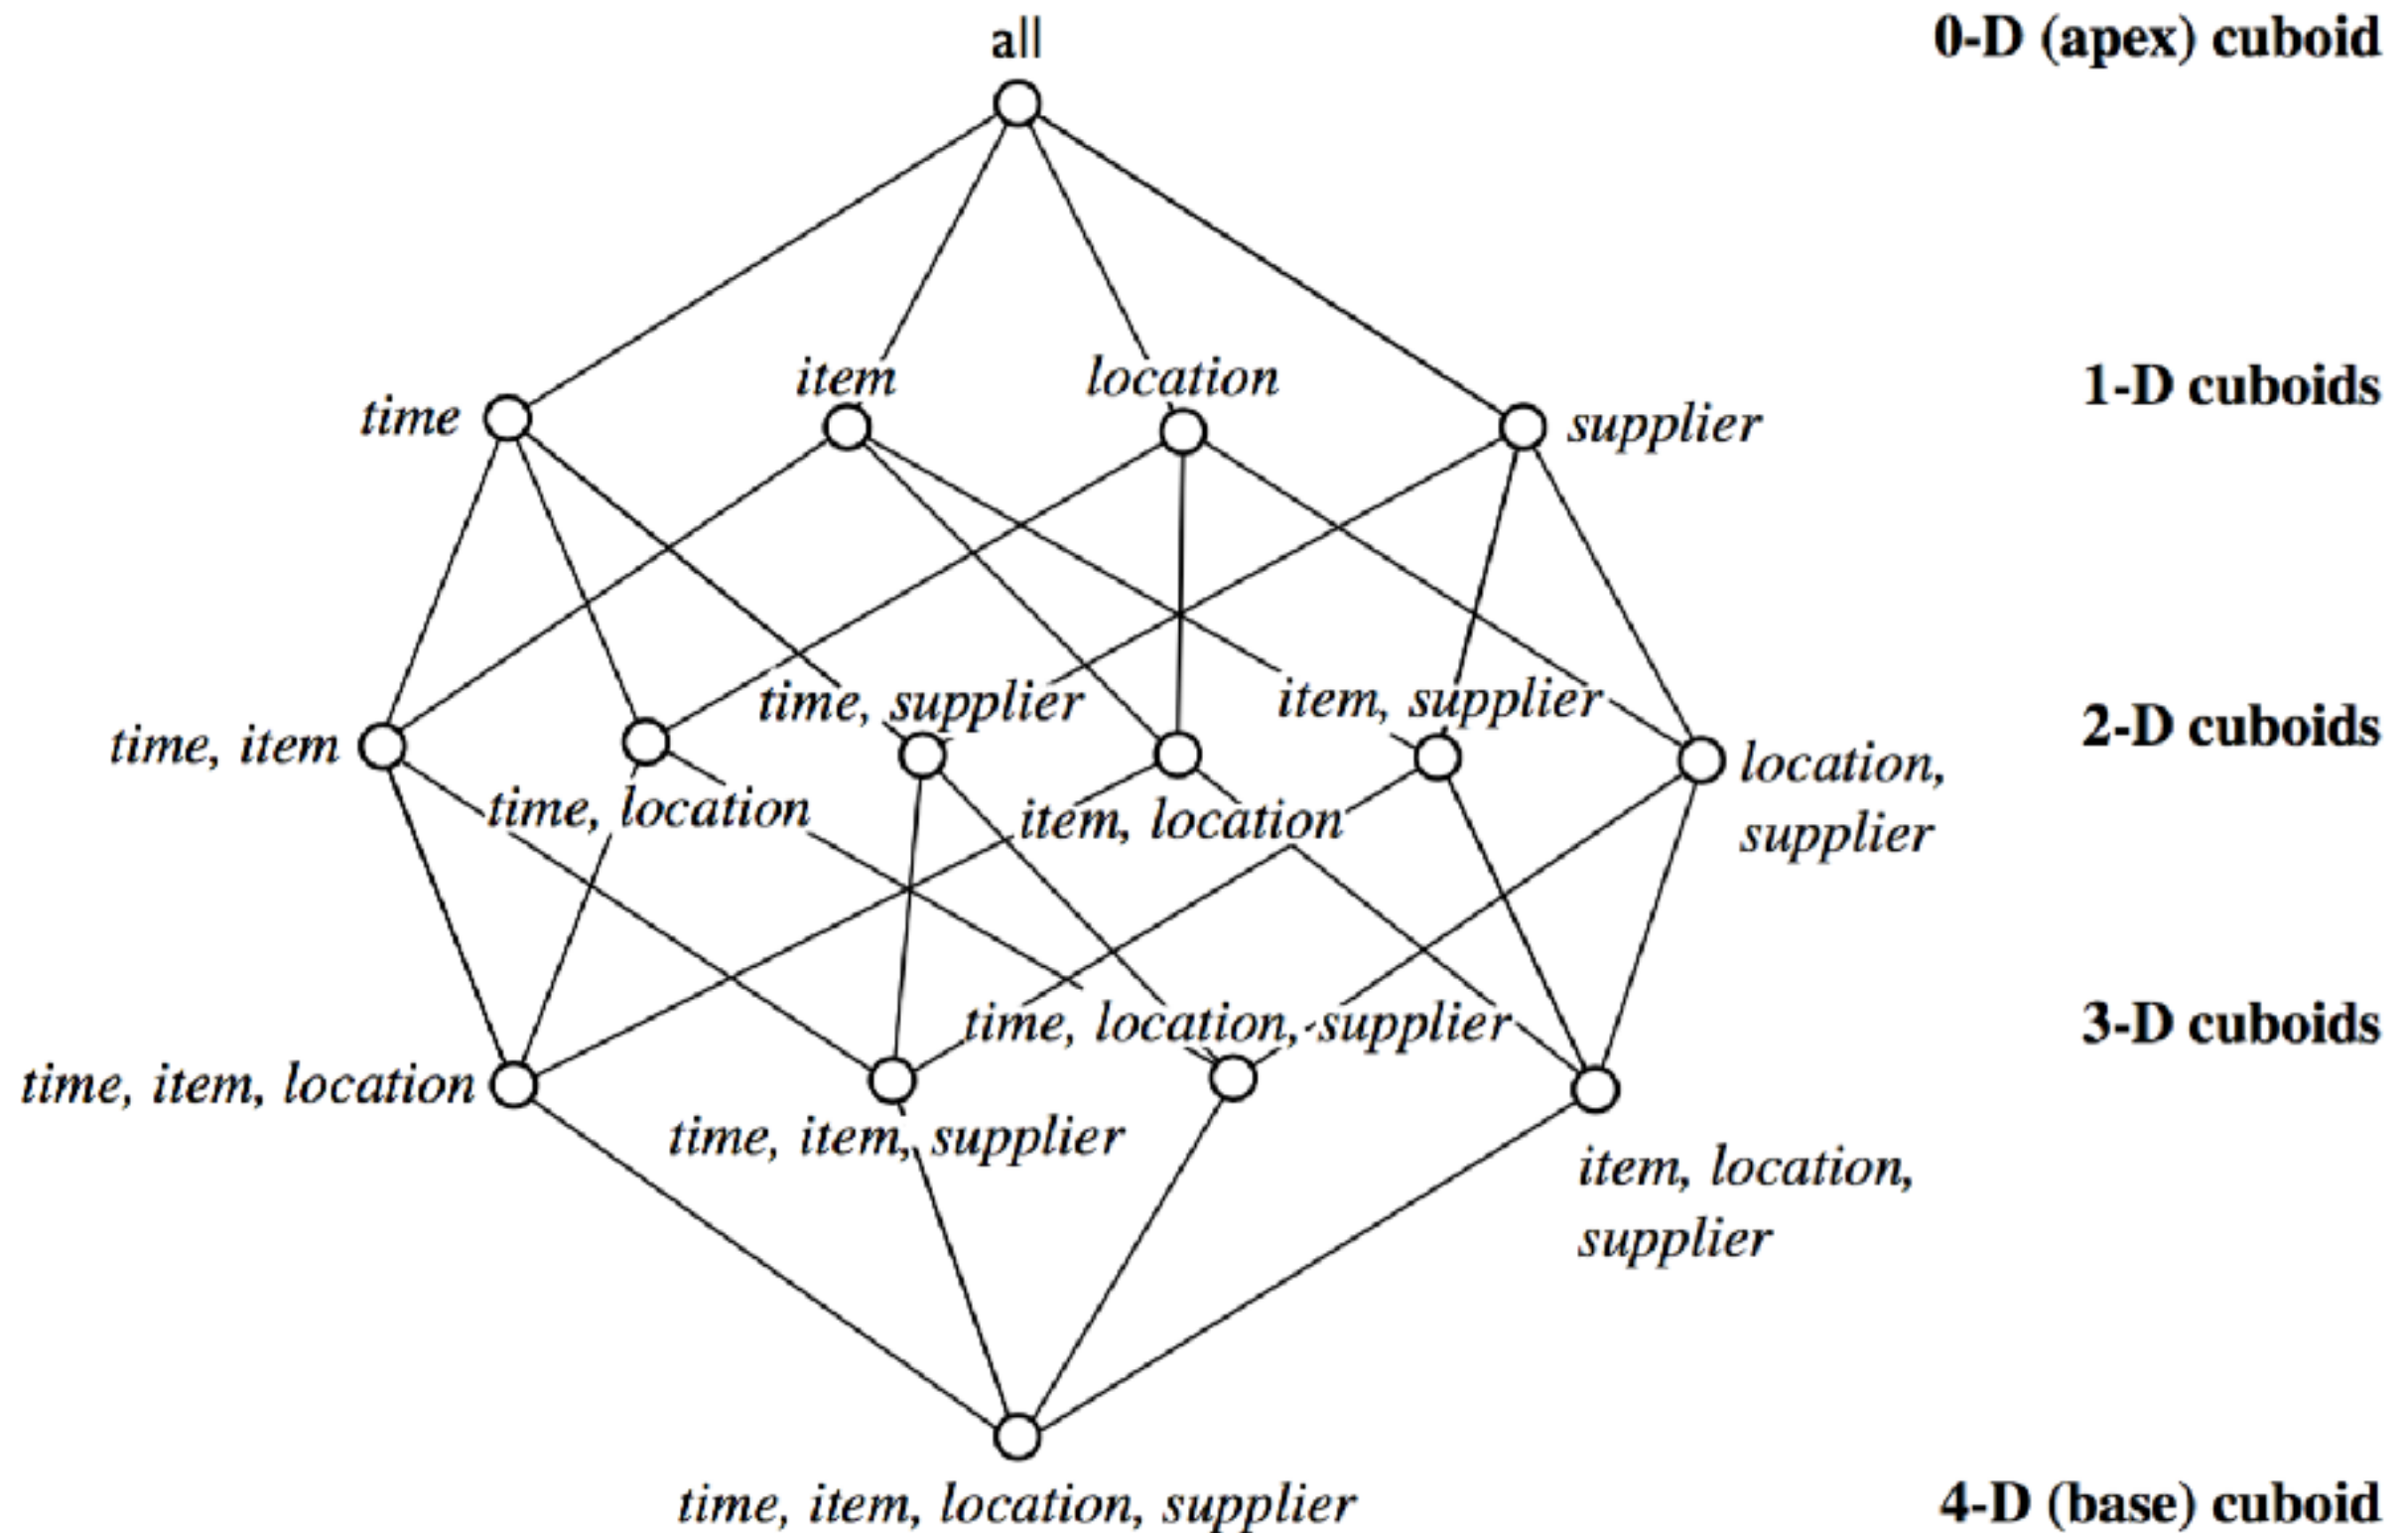
zero tolerance policy!

You are encouraged to form a **study group** to discuss the homework and the programming assignments but are expected to complete the homework and programming assignments **completely on your own without recourse to notes from the group discussions.**

Plagiarism: It is an academic violation to copy, to include text from other sources, including online sources, without proper citation.

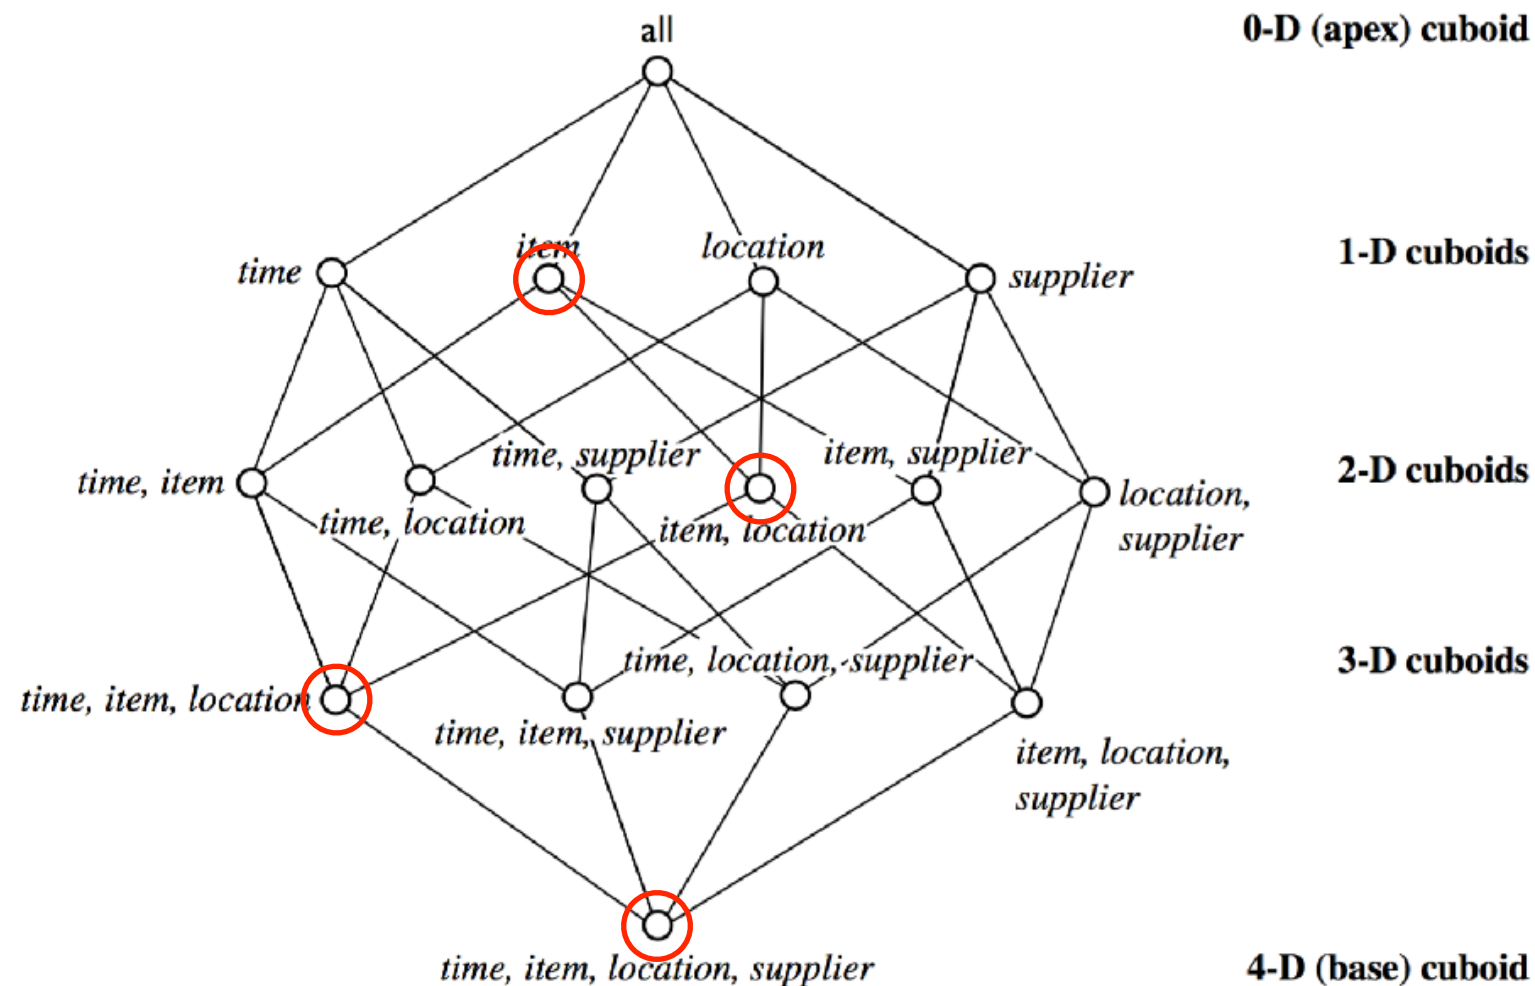
Any student found to
be **violating** this code
will be subject to
disciplinary action.

data cube: a lattice of cuboids



DATA CUBE: A LATTICE OF CUBOIDS

Base vs. aggregate cells;
ancestor vs. descendant cells;
parent vs. child cells



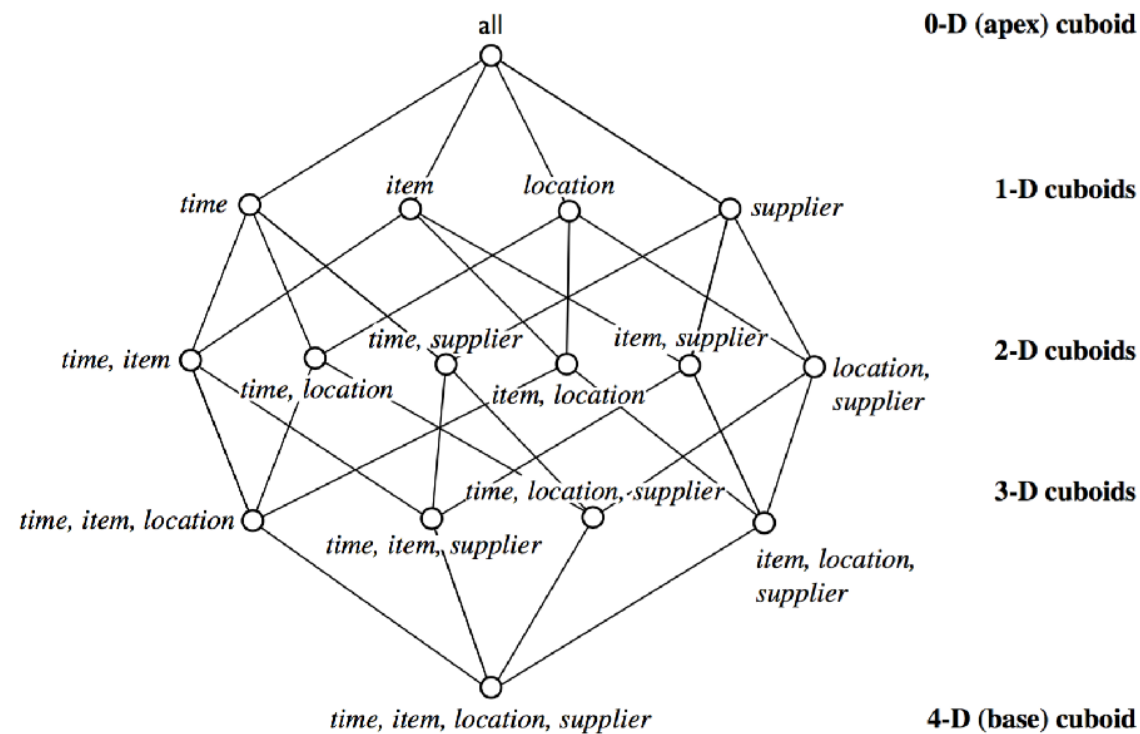
(9/15, milk, Urbana,
Dairy_land)

(9/15, milk, Urbana, *)

(*, milk, Urbana, *)

(*, milk, Chicago, *)

(*, milk, *, *)



CUBE MATERIALIZATION

Full cube vs. iceberg cube

compute cube sales iceberg as

select month, city, customer
group, count(*)


from salesInfo

cube by month, city, customer
group

having **count(*) >= min support**

Computing only the cuboid
cells whose measure satisfies
the iceberg condition

Only a small portion of cells
may be “above the water” in a
sparse cube



a priori principle

How many aggregate cells
if "having count >= 1"?

	Q1	Q2	Q3	Q4
Milk	5	1	4	1
Poultry	1	10	5	3
Beer	6	7	1	11
Chocolate	1	8	12	2

base cell

	Q1	Q2	Q3	Q4
Milk	5	1	4	1
Poultry	1	10	5	3
Beer	6	7	1	11
Chocolate	1	8	12	2

base cell

how many aggregates?

(Milk, Q1)

(Beer, Q2)

(Milk,*) (Beer,*)

(*,*)

(*,Q1) (*,Q2)

$$3 + 3 - 1 = 5$$

Exercise!

TIME	USA	CAN	FRA	ITL
1997	300	845	785	129
1998	139	481	782	
1999	974	121	312	194
	328	465	513	120
	545	741	962	51
	745	159	901	

Legend: Total Sales, Expenses

$\{ (a_1, a_2, a_3, \dots, a_{100}): 10, (b_1, b_2, b_3, \dots, b_{100}): 10 \}$

how many aggregate cells have more than or equal to 10?

$$2^{100} + 2^{100} - 1 - 2 = 2^{101} - 3$$

one common aggregate cell (*)

two base cells

Exercise!

		PRODUCT			
		Product A		Product B	
TIME	1997	300	845	785	129
		139	481	782	
	1998	974	121	312	194
		328	465	513	120
	1999	545	741	962	51
		745	159	901	
		USA	CAN	FRA	ITL
		North America		Europe	
		LOCATION			

Total Sales
 Expenses


2^2

 $\{ (\underline{a_1}, a_2, a_3, \dots, a_{100}) : 10, (\underline{a_1}, a_2, b_3, \dots, b_{100}) : 10 \}$

how many aggregate cells have counts greater than or equal to 10?

$$2^{100} + 2^{100} - 4 - 2 = 2^{101} - 6$$

four common aggregate cells
 two base cells



Closed cell c : if there exists no cell d , such that d is a descendant of c , and d has the same measure value as c .

	Q1	Q2	Q3	Q4
Milk	5	1	4	1
Poultry	1	10	5	3
Beer	6	7	1	11
Chocolate	1	8	12	2

closed cube

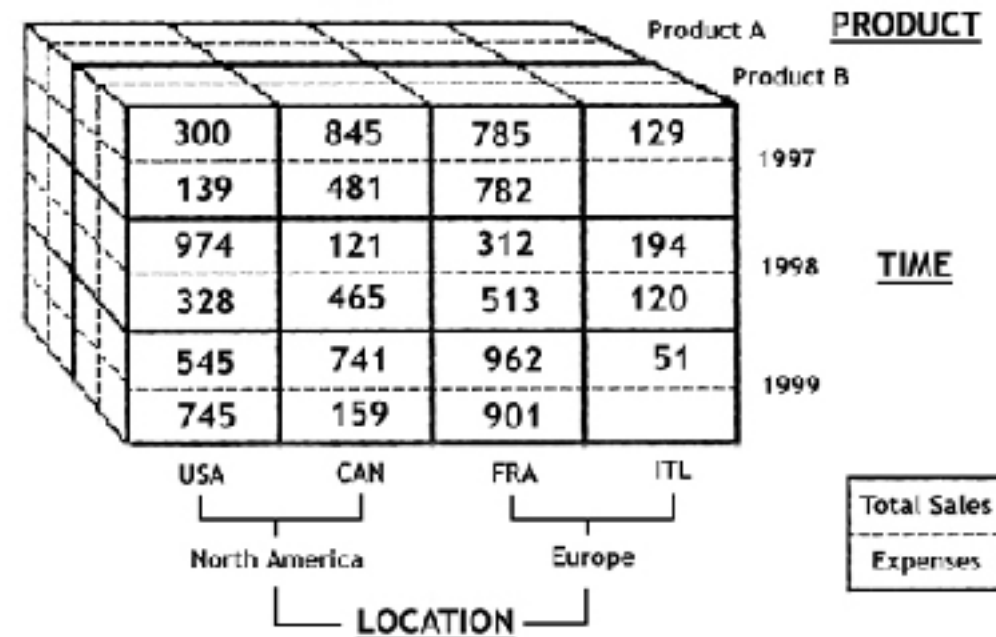
Exercise!

		<u>PRODUCT</u>				
		Product A		Product B		
	1997	300	845	785	129	
	1998	139	481	782		
	1999	974	121	312	194	
		328	465	513	120	
		545	741	962	51	
		745	159	901		
		USA	CAN	FRA	ITL	
		North America		Europe		
		<u>LOCATION</u>				

Total Sales
 Expenses

$\{ (a_1, a_2, a_3, \dots, a_{100}): 10, (a_1, a_2, b_3, \dots, b_{100}): 10 \}$

how many closed cells?



cube shell

Pre-compute only the cuboids involving a small # of dimensions, e.g., 3

More dimension combinations will need to be computed on the fly

EFFICIENT COMPUTATION ROADMAP

.....

General cube computation heuristics
(Agarwal et al.'96)

Computing full/iceberg cubes: 3
methodologies

Bottom-Up: Multi-Way array aggregation
(Zhao, Deshpande & Naughton, SIGMOD'97)

Top-down:

BUC (Beyer & Ramakrishnan, SIGMOD'99)

H-cubing technique (Han, Pei, Dong & Wang:
SIGMOD'01)

Integrating Top-Down and Bottom-Up:


Star-cubing algorithm (Xin, Han, Li & Wah:
VLDB'03)

High-dimensional OLAP: A Minimal Cubing
Approach (Li, et al. VLDB'04)

Computing alternative kinds of cubes:

Partial cube, closed cube, approximate cube,
etc.

GENERAL HEURISTICS



Agarwal, S., Agrawal, R., Deshpande, P. M., Gupta, A., Naughton, J. F., Ramakrishnan, R., & Sarawagi, S. (1996, September). **On the computation of multidimensional aggregates**. In VLDB (Vol. 96, pp. 506-521).

Sorting, hashing, and grouping operations are applied to the dimension attributes in order to reorder and cluster related tuples

Aggregates may be computed from previously computed aggregates, rather than from the base fact table

Smallest-child: computing a cuboid from the smallest, previously computed cuboid

Cache-results: caching results of a cuboid from which other cuboids are computed to reduce disk I/Os

Amortize-scans: computing as many as possible cuboids at the same time to amortize disk reads


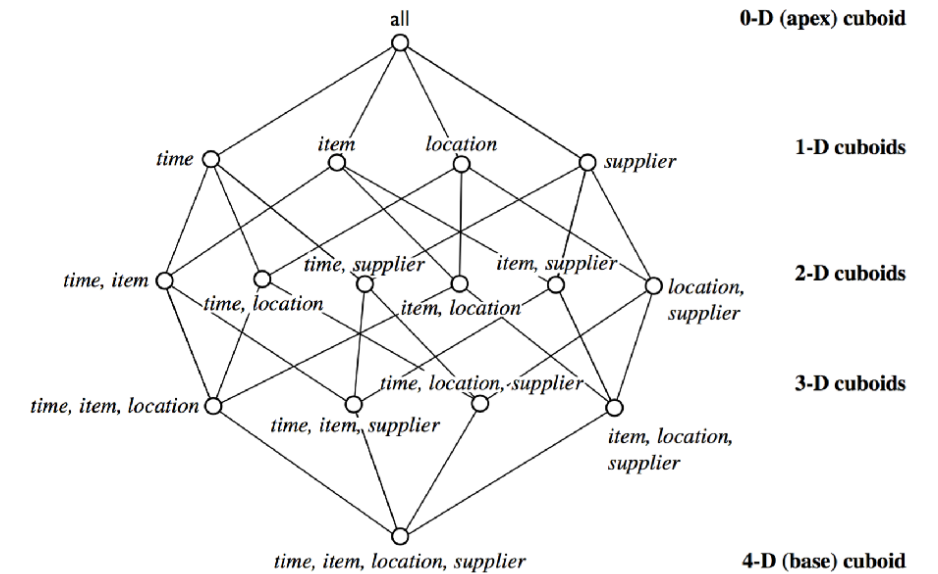
Share-sorts: sharing sorting costs across multiple cuboids when sort-based method is used

Share-partitions: sharing the partitioning cost across multiple cuboids when hash-based algorithms are used

Closed cell **c**: if there exists no cell **d**, such that **d** is a descendant of **c**, and **d** has the same measure value as **c**.

BASIC CONCEPTS SUMMARY

Methods Advanced Processing Multidimensional Summary



iceberg cube

$$2^n - 1$$

base cell aggregates

Multi-Way Array Aggregation

BUC

High-Dimensional OLAP

DATA CUBE METHODS

.....

Basic Concepts Advanced Processing Multidimensional Summary

MULTI-WAY AGGREGATION

Array-based “bottom-up” algorithm


Using **multi-dimensional** chunks

No direct tuple comparisons

Simultaneous aggregation on multiple dimensions

Intermediate aggregate values are **re-used** for computing ancestor cuboids

Cannot do *Apriori* pruning:
No iceberg optimization



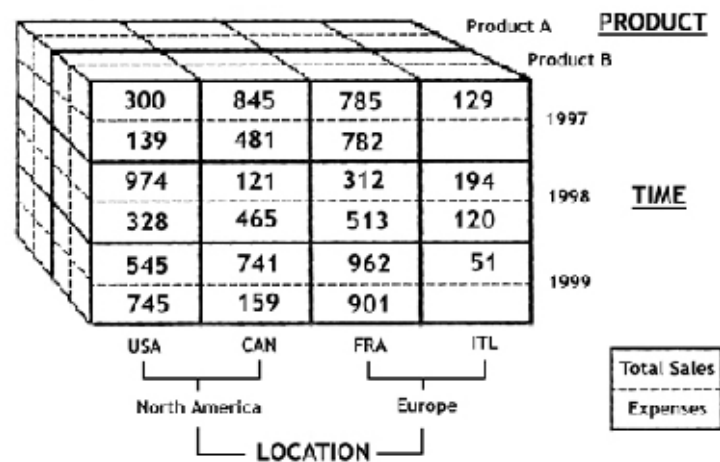
MOLAP

.....


Partition arrays into chunks
(a small subcube which fits in
memory).

Compressed sparse array
addressing: (chunk_id, offset)

Compute aggregates in
“multiway” by visiting cube
cells in the order which
minimizes the number of cell
visits, and which reduces
memory access and storage
cost.




aggregating up one dimension



Can we avoid revisiting cells
to compute different
aggregates?





what happens when we re-order our calculations?



smallest plane

2nd smallest plane



single
dimensional
aggregates

MOLAP SUMMARY


.....

Method: the planes should be sorted and computed according to their size in ascending order

Idea: keep the smallest plane in the main memory, fetch and compute only one chunk at a time for the largest plane

Limitation of the method: performs well only for a small number of dimensions

If there are a large number of dimensions, “top-down” computation and iceberg cube computation methods can be explored




data is not sparse
moderate dimensionality

failure cases

apriori

Bottom Up Construction



BOTTOM UP CONSTRUCTION

.....

Bottom-up cube computation

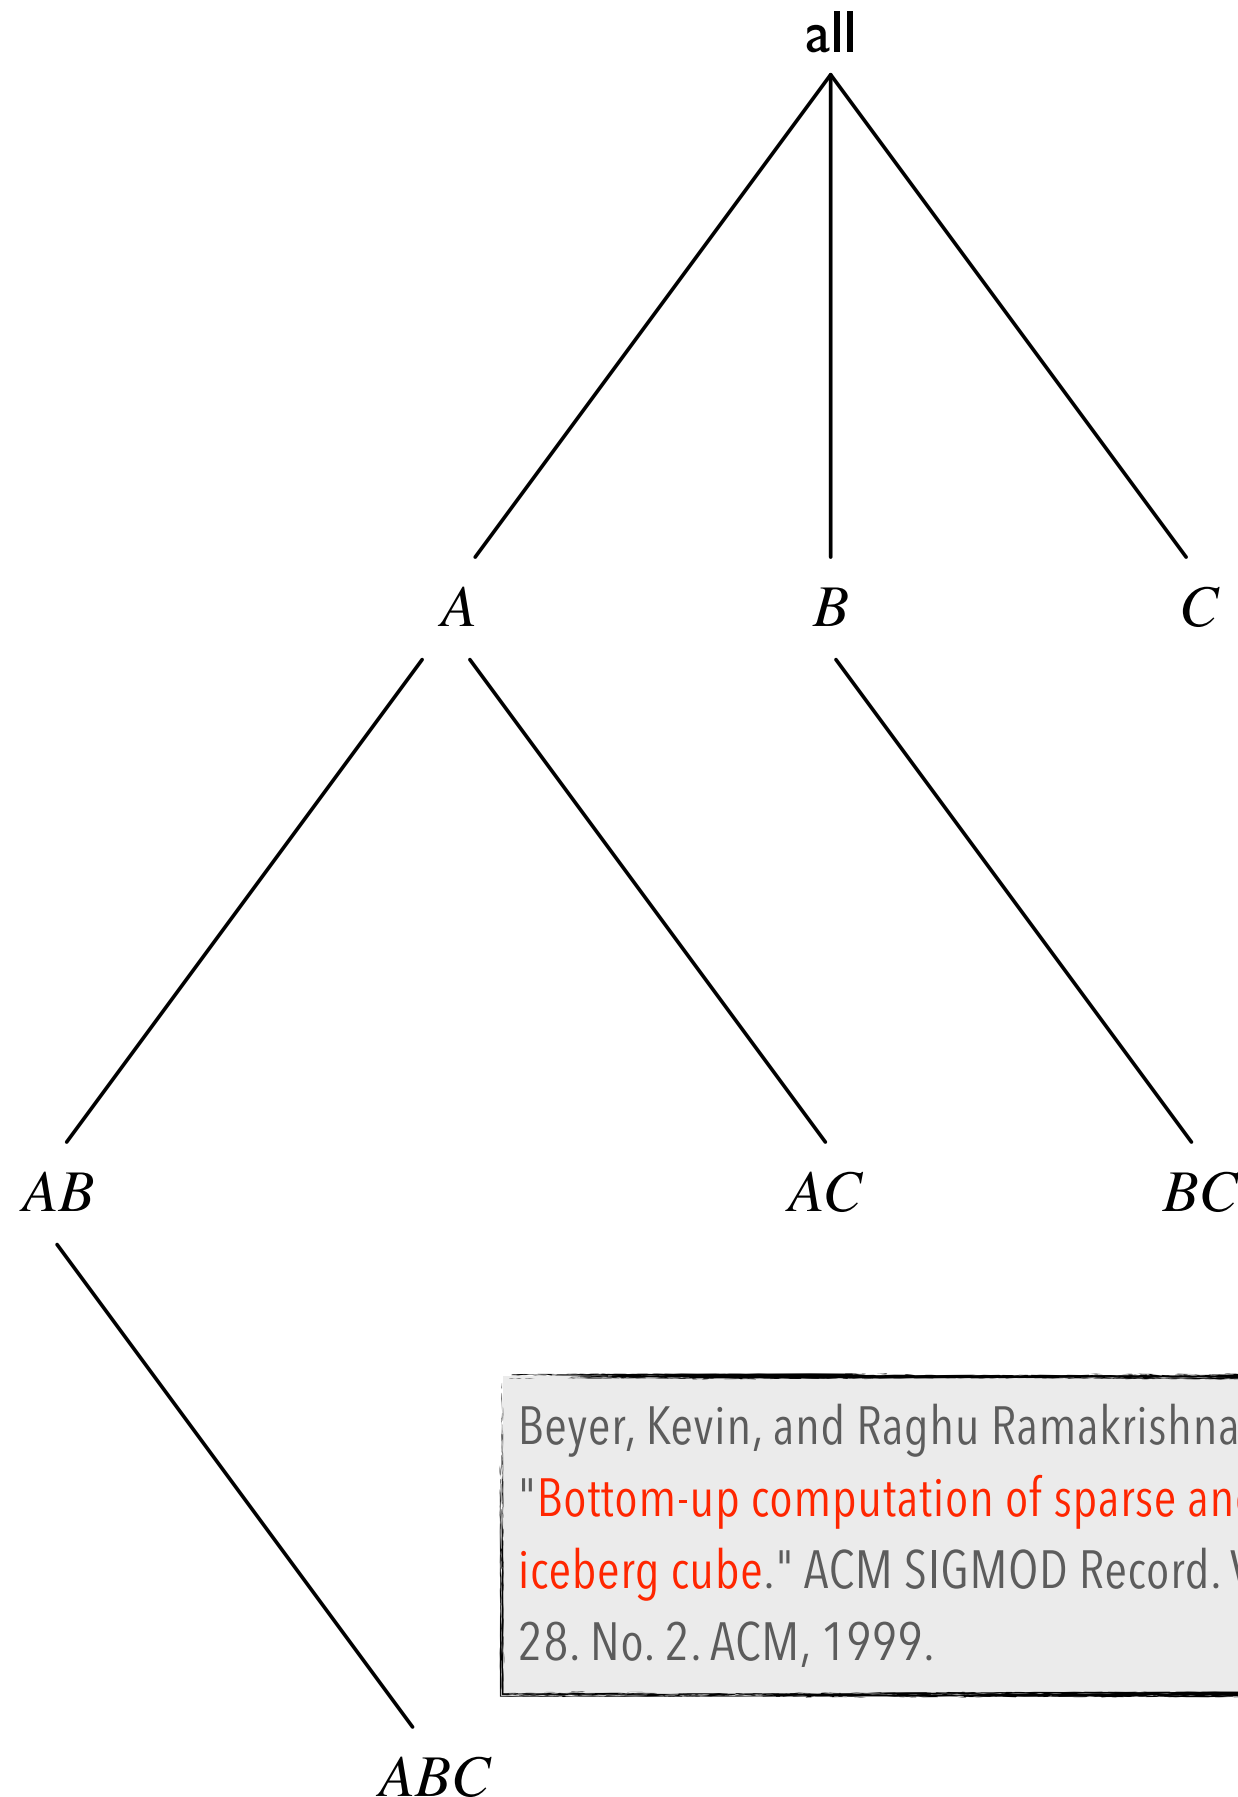
(Note: top-down in our view!)

Divides dimensions into partitions and facilitates iceberg pruning


If a partition does not satisfy min_sup , its descendants can be pruned

If $\text{min_sup} = 1 \Rightarrow$ compute full CUBE!

No simultaneous aggregation



Beyer, Kevin, and Raghu Ramakrishnan.
"Bottom-up computation of sparse and
iceberg cube." ACM SIGMOD Record. Vol.
28. No. 2. ACM, 1999.



BUC: PARTITIONING

Usually, entire data set
can't fit in main memory

Sort distinct values

partition into blocks that fit

Continue processing

Optimizations

Partitioning

External Sorting, Hashing,
Counting Sort

Ordering dimensions to
encourage pruning

Cardinality, Skew,
Correlation

COUNT SORT

.....

For simplicity, consider the data in the range 0 to 9.

Input data: 1, 4, 1, 2, 7, 5, 2

1) Take a count array to store the count of each unique object.

Index:	0	1	2	3	4	5	6	7	8	9
Count:	0	2	2	0	1	1	0	1	0	0

2) Modify the count array such that each element at each index stores the sum of previous counts.

Index:	0	1	2	3	4	5	6	7	8	9
Count:	0	2	4	4	5	6	6	7	7	7

The modified count array indicates the position of each object in the output sequence.

3) Output each object from the input sequence followed by

Decreasing its count by 1.

Process the input data: 1, 4, 1, 2, 7, 5, 2. Position of 1 is 2.

Put data 1 at index 2 in output. Decrease count by 1 to place the next data value of 1 at an index one smaller than this index.

A:4, B:4, C:2, D:2

a_1	b_1	c_1					
		c_2					
	b_2						
	b_3						
	b_4						
a_2 CountSort for each dimension							
a_3							
a_4							

BUC: OVERVIEW

Start from partition = Apex
 (*, *, *, *)

Have an order of dimensions

Count current partition

Sort **distinct** values in next dimension

Partition into distinct values

Recurse into each partition

Algorithm: BUC. Algorithm for the computation of sparse and iceberg cubes.

Input:

- *input*: the relation to aggregate;
- *dim*: the starting dimension for this iteration.


Globals:

- constant *numDims*: the total number of dimensions;
- constant *cardinality[numDims]*: the cardinality of each dimension;
- constant *min_sup*: the minimum number of tuples in a partition for it to be output;
- *outputRec*: the current output record;
- *dataCount[numDims]*: stores the size of each partition. *dataCount[i]* is a list of integers of size *cardinality[i]*.

Output: Recursively output the iceberg cube cells satisfying the minimum support.

Method:


- (1) Aggregate(input); // Scan *input* to compute measure, e.g., count. Place result in *outputRec*.
- (2) **if** input.count() == 1 **then** // Optimization
 WriteDescendants(input[0], dim); **return**;
 endif
- (3) write outputRec;
- (4) **for** ($d = \text{dim}$; $d < \text{numDims}$; $d++$) **do** //Partition each dimension
- (5) $C = \text{cardinality}[d]$;
- (6) Partition(input, d, C, dataCount[d]); //create C partitions of data for dimension d
- (7) $k = 0$;
- (8) **for** ($i = 0$; $i < C$; $i++$) **do** // for each partition (each value of dimension d)
- (9) $c = \text{dataCount}[d][i]$;
- (10) **if** $c \geq \text{min_sup}$ **then** // test the iceberg condition
- (11) outputRec.dim[d] = input[k].dim[d];
- (12) BUC(input[k..k + c - 1], $d + 1$); // aggregate on next dimension
- (13) **endif**
- (14) $k += c$;
- (15) **endfor**
- (16) outputRec.dim[d] = all;
- (17) **endfor**




```

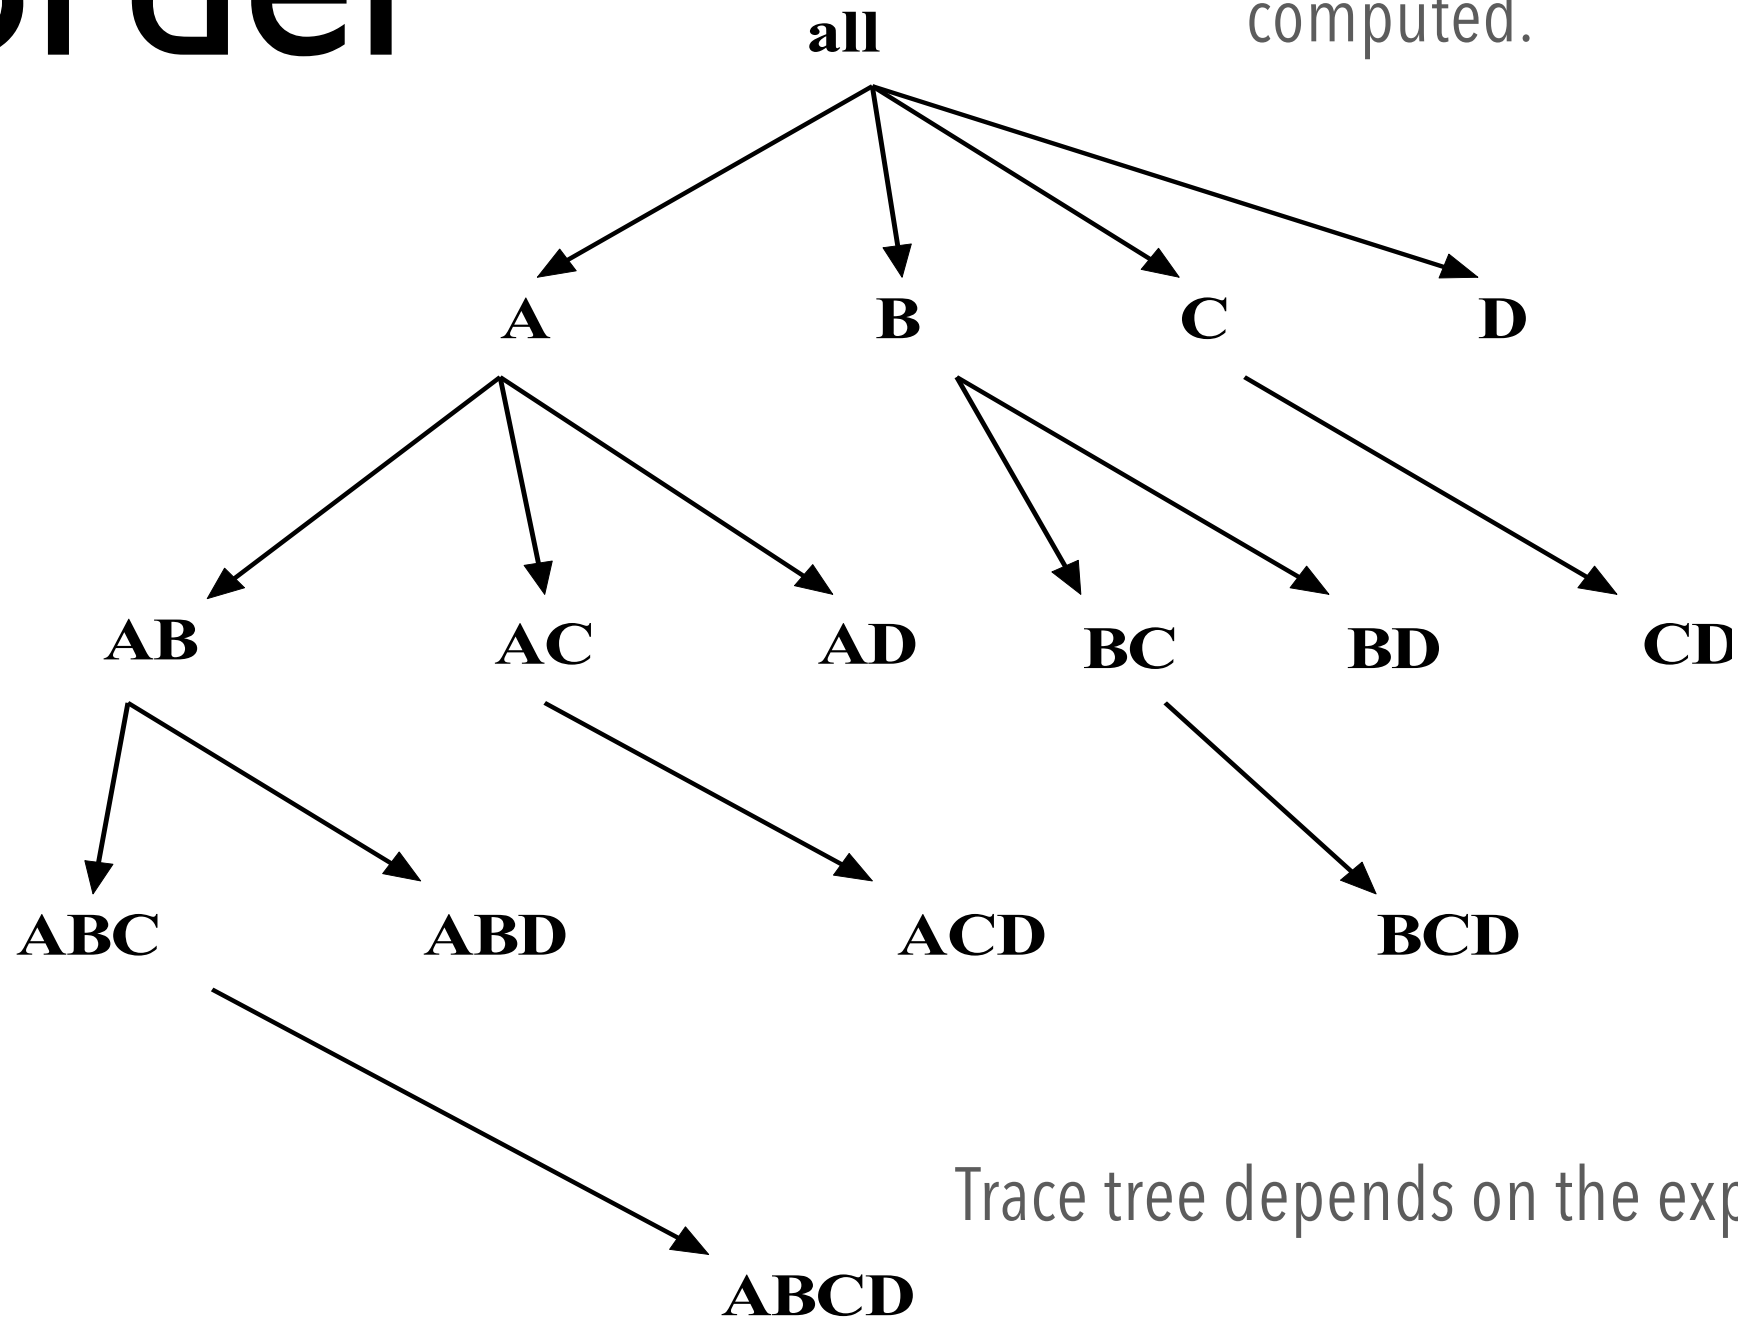
(1) Aggregate(input); // Scan input to compute measure, e.g., count. Place result in outputRec.
(2) if input.count() == 1 then // Optimization
    WriteDescendants(input[0], dim); return;
endif
(3) write outputRec;
(4) for ( $d = \text{dim}; d < \text{numDims}; d++$ ) do // Partition each dimension
(5)      $C = \text{cardinality}[d];$  ← the total number of dimensions
(6)     Partition(input, d, C, dataCount[d]); //create C partitions of data for dimension  $d$ 
(7)      $k = 0;$ 
(8)     for ( $i = 0; i < C; i++$ ) do // for each partition (each value of dimension  $d$ )
(9)          $c = \text{dataCount}[d][i];$  ← the number of times a specific attribute value occurs
(10)        if  $c \geq \text{min\_sup}$  then // test the iceberg condition
(11)            outputRec.dim[d] = input[k].dim[d];
(12)            BUC(input[k..k + c - 1],  $d + 1$ ); // aggregate on next dimension
(13)        endif
(14)         $k += c;$  ← recursively call BUC
(15)    endfor
(16)    outputRec.dim[d] = all;
(17) endfor

```



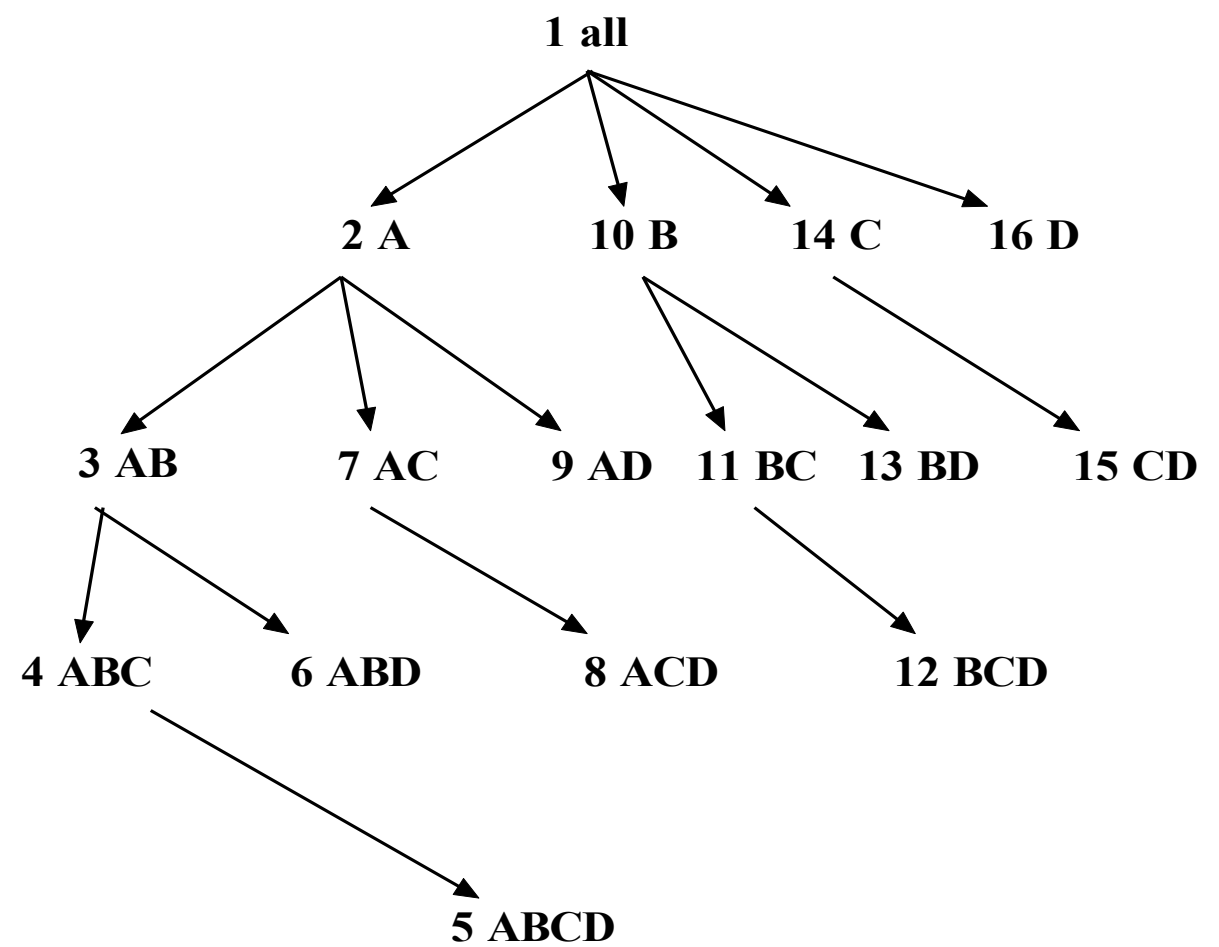

exploration order


Different orders may
result in different
numbers of cells to be
computed.



Trace tree depends on the exploration order.

$A \rightarrow B \rightarrow C \rightarrow D$





skew

cardinality


order

sensitivities

no aggregate sharing

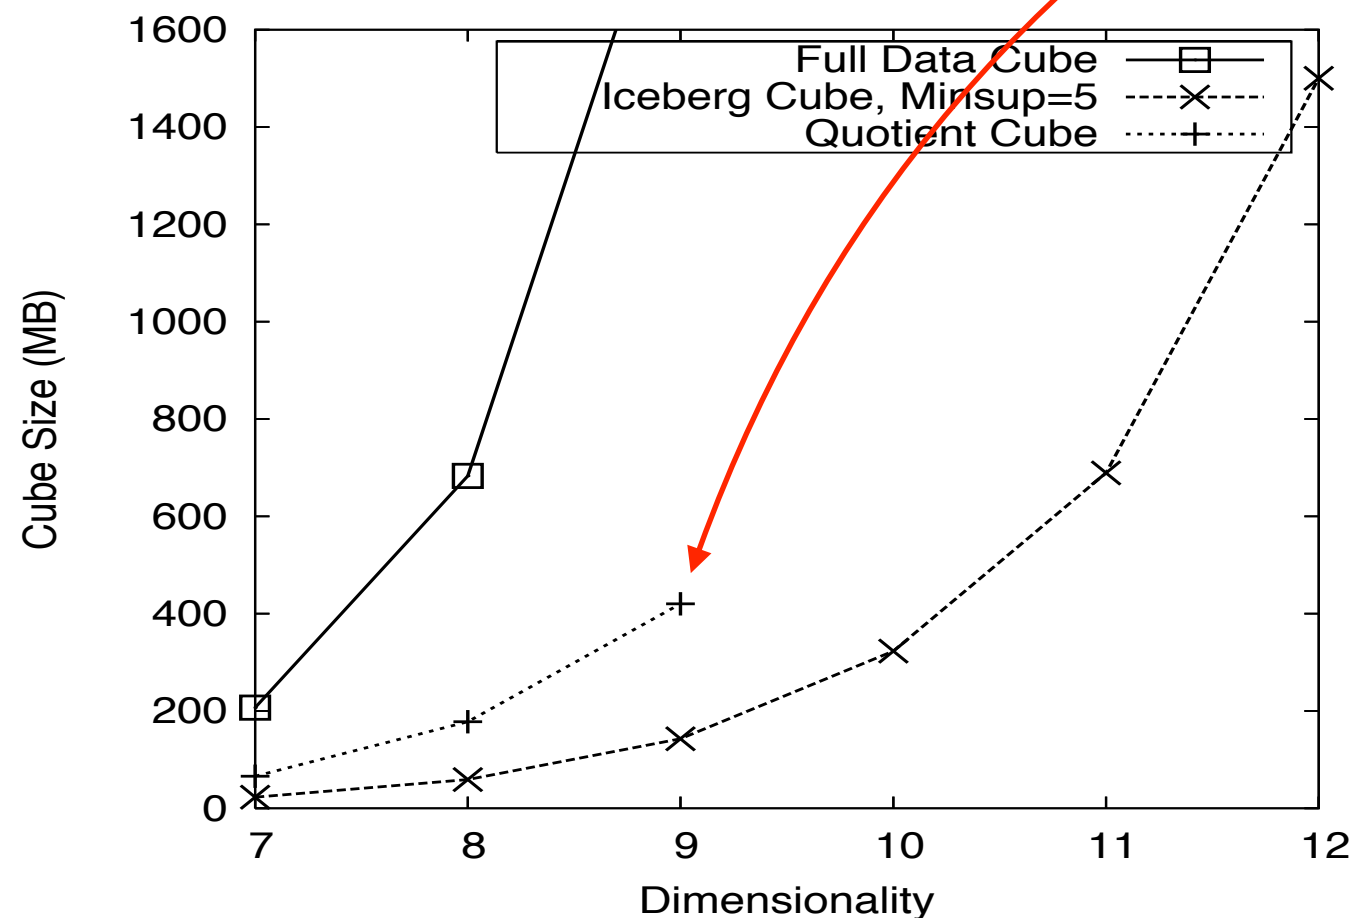
what about high
dimensions?





Li, Xiaolei, Jiawei Han, and Hector Gonzalez. "High-dimensional OLAP: a minimal cubing approach." Proceedings of the Thirtieth international conference on Very large data bases-Volume 30. VLDB Endowment, 2004.

ICEBERG CHALLENGES



First, if a high-dimensional cell has the support already passing the iceberg threshold, **it cannot be pruned** by the iceberg condition and will still generate a huge number of cells. For example, a base- cuboid cell:“(a₁, a₂, ..., a₆₀): 5” (i.e., with count 5) will still generate 2⁶⁰ iceberg cube cells.

Second, it is **difficult** to set up an appropriate iceberg threshold. A too low threshold will still generate a huge cube, but a too high one may invalidate many useful applications.

Third, an iceberg cube **cannot be incrementally updated**. Once an aggregate cell falls below the iceberg threshold and is pruned, incremental update will not be able to recover the original measure.



CHALLENGES

The “curse of dimensionality” problem

Iceberg cube and compressed cubes: only delay the inevitable explosion

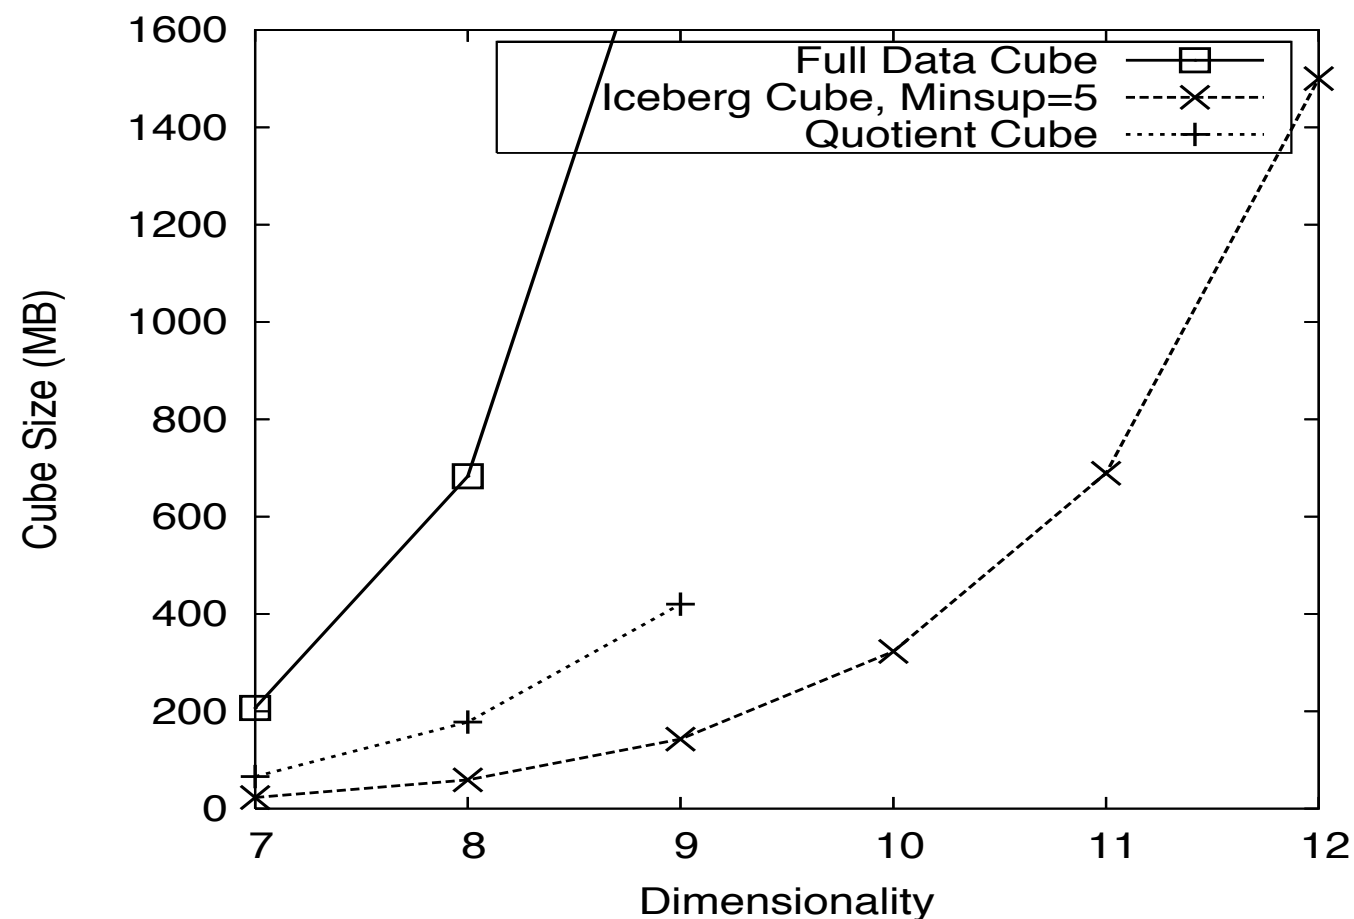
Full materialization: still significant overhead in accessing results on disk

High-D OLAP is needed in applications

Science and engineering analysis


Bio-data analysis: thousands of genes

Statistical surveys: hundreds of variables



Observation: OLAP occurs only on a small subset of dimensions at a time

Semi-Online
Computational Model



High Dimensional OLAP

Partition the set of dimensions into shell fragments

Compute data cubes for each shell fragment while retaining **inverted** indices or value-list indices

Given the **pre-computed** fragment cubes, **dynamically** compute cube cells of the high-dimensional data cube **online**

Partitions the data vertically

Reduces high-
dimensional cube into
a set of lower
dimensional cubes


properties

Online re-
construction of
original high-
dimensional space

Lossless reduction

Offers tradeoffs between
the amount of pre-
processing and the speed
of online computation

Divide the 5-D table into 2 shell fragments: (A, B, C) and (D, E)




<i>tid</i>	A	B	C	D	E
1	a1	b1	c1	d1	e1
2	a1	b2	c1	d2	e1
3	a1	b2	c1	d1	e2
4	a2	b1	c1	d1	e2
5	a2	b1	c1	d1	e3

measure: count

The inverted index table uses the **same amount** of storage space as the original database.

<i>tid</i>	A	B	C	D	E
1	a1	b1	c1	d1	e1
2	a1	b2	c1	d2	e1
3	a1	b2	c1	d1	e2
4	a2	b1	c1	d1	e2
5	a2	b1	c1	d1	e3



Attribute Value	TID List	List Size
a1	1 2 3	3
a2	4 5	2
b1	1 4 5	3
b2	2 3	2
c1	1 2 3 4 5	5
d1	1 3 4 5	4
d2	2	1
e1	1 2	2
e2	3 4	2
e3	5	1

Build traditional inverted index or RID list

Attribute Value	TID List	List Size
a1	1 2 3	3
a2	4 5	2
b1	1 4 5	3
b2	2 3	2
c1	1 2 3 4 5	5
d1	1 3 4 5	4
d2	2	1
e1	1 2	2
e2	3 4	2
e3	5	1

<i>Cell</i>	<i>Intersection</i>	<i>TID List</i>	<i>List Size</i>
(a_1, b_1)	$\{1, 2, 3\} \cap \{1, 4, 5\}$	$\{1\}$	1
(a_1, b_2)	$\{1, 2, 3\} \cap \{2, 3\}$	$\{2, 3\}$	2
(a_2, b_1)	$\{4, 5\} \cap \{1, 4, 5\}$	$\{4, 5\}$	2
(a_2, b_2)	$\{4, 5\} \cap \{2, 3\}$	$\{\}$	0

SHELL FRAGMENTS

.....

Generalize the I-D inverted indices to multi-dimensional ones in the data cube sense

Compute all cuboids for data cubes ABC and DE while retaining the inverted indices

For example, shell fragment cube ABC contains 7 cuboids:

A, B, C

AB, AC, BC

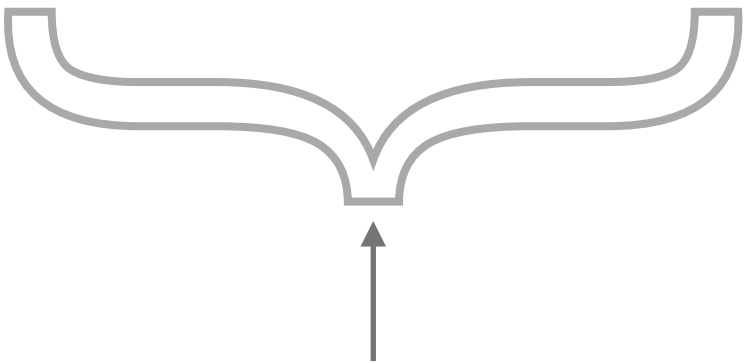
ABC

This completes the offline computation stage

SHELL SIZE AND DESIGN

.....

Given a database of T tuples, D dimensions, and F shell fragment size, the fragment cubes' space requirement is:

$$O \left(T \left[\frac{D}{F} \right] (2^F - 1) \right)$$


each tuple id will be stored in these many cuboids

Shell fragments do not have to be disjoint

Fragment groupings can be **arbitrary** to allow for maximum online performance

Known common combinations (e.g., <city, state>) should be grouped together.

Shell fragment sizes can be adjusted for optimal balance between offline and online computation

If measures other than **count** are present, store in **ID_measure** table separate from the shell fragments

tid	count	sum
1	5	70
2	3	10
3	8	20
4	5	40
5	2	30

FRAG SHELLS ALGORITHM

.....

Partition set of dimension (A_1, \dots, A_n) into a set of k fragments (P_1, \dots, P_k) .

Scan base table once and do the following:

insert $\langle \text{tid}, \text{measure} \rangle$ into ID_measure table.

for each attribute value a_i of each dimension A_i

build inverted index entry $\langle a_i, \text{tidlist} \rangle$

For each fragment partition P_i

build local fragment cube S_i by intersecting tid-lists in bottom-up fashion.

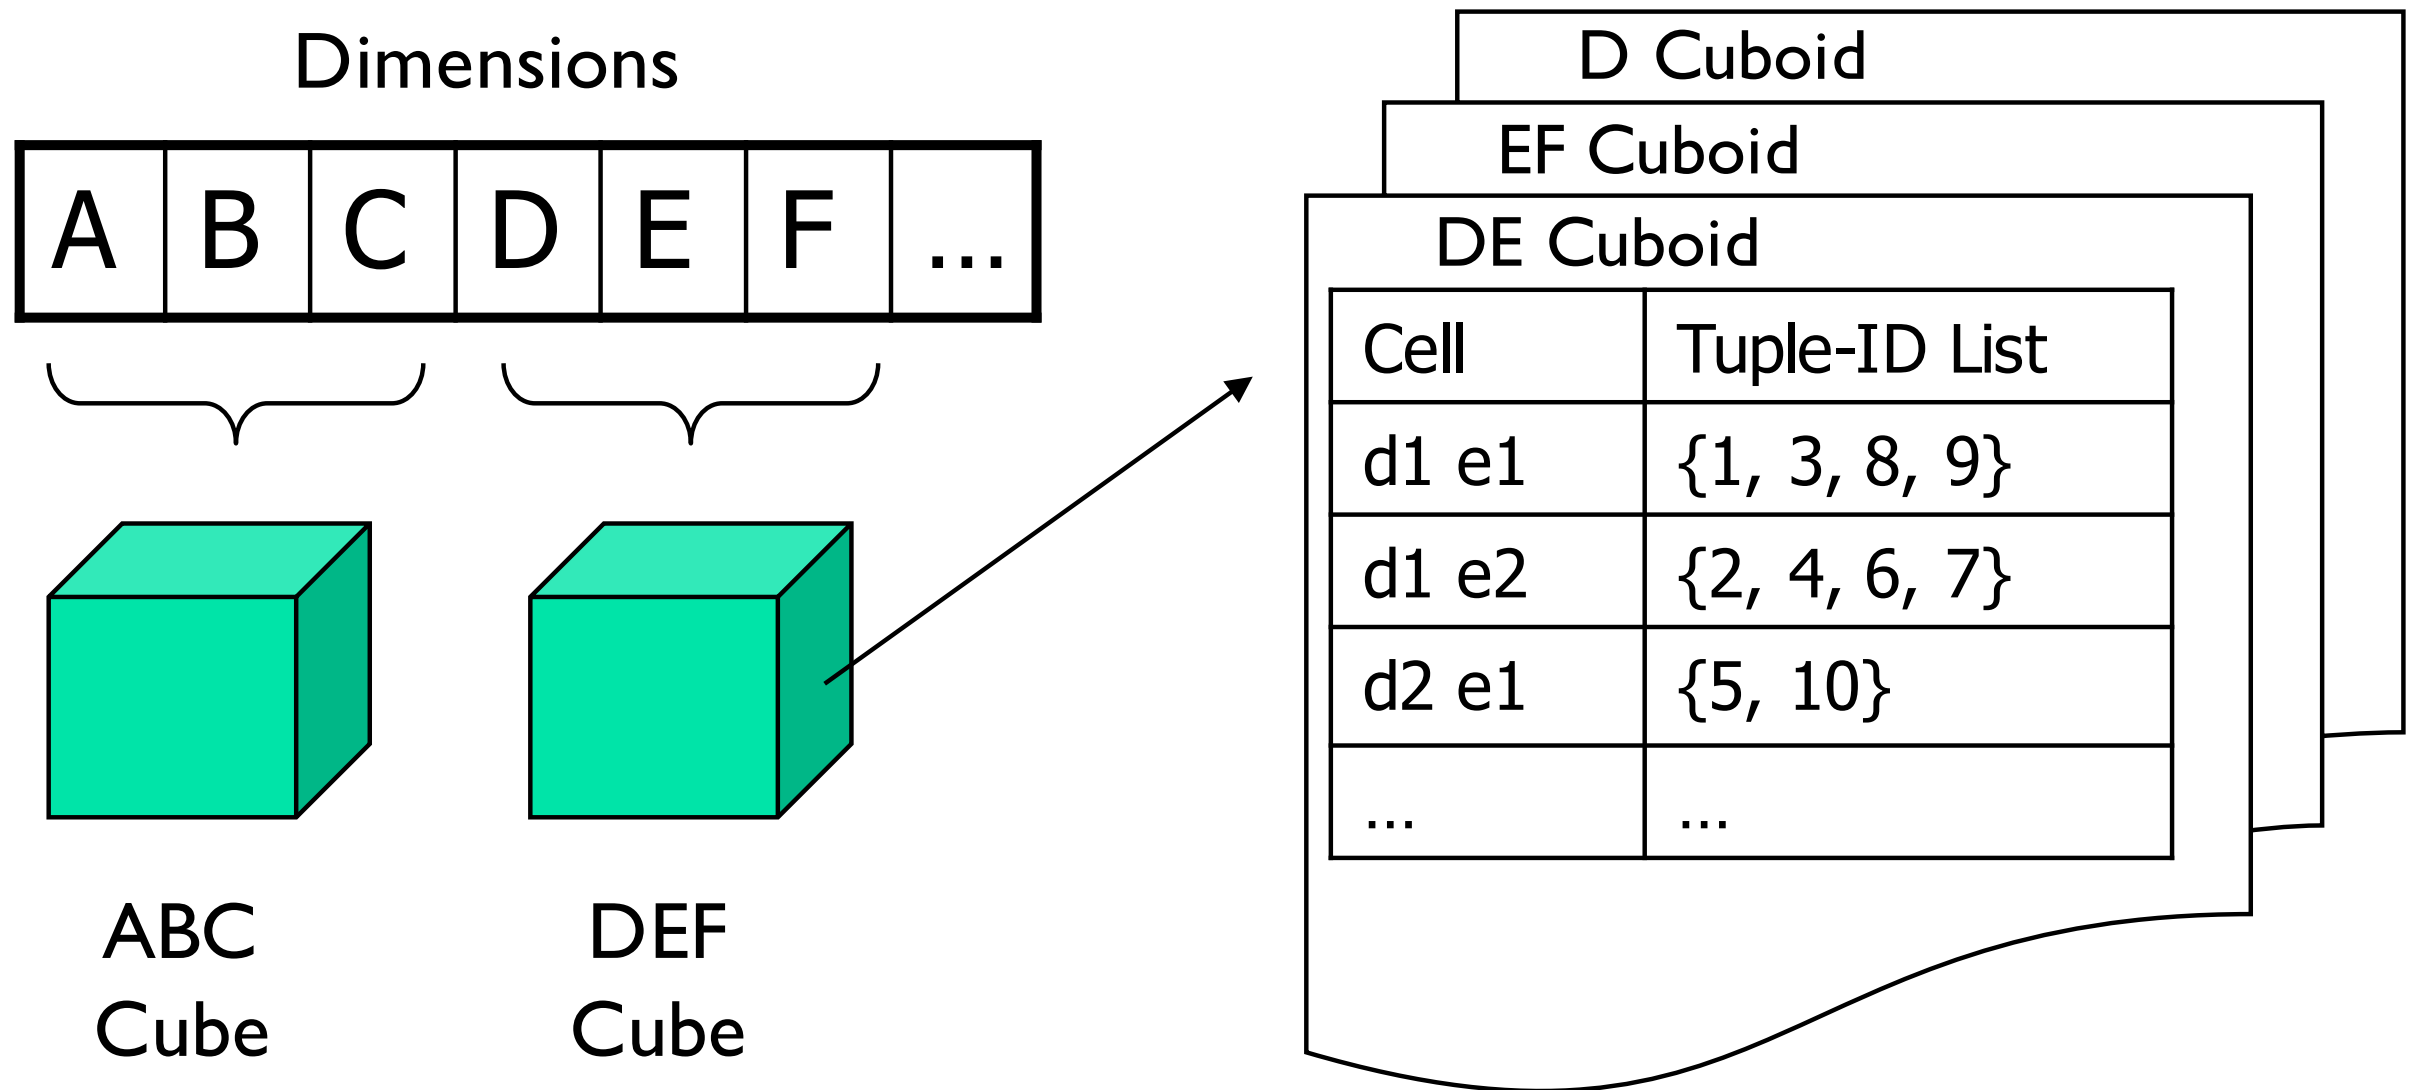
measure array

tid	count	sum
1	5	70
2	3	10
3	8	20
4	5	40
5	2	30

Attribute Value	TID List	List Size
a1	1 2 3	3
a2	4 5	2
b1	1 4 5	3
b2	2 3	2
c1	1 2 3 4 5	5
d1	1 3 4 5	4
d2	2	1
e1	1 2	2
e2	3 4	2
e3	5	1

Cell	Intersection	TID List	List Size
(a_1, b_1)	$\{1, 2, 3\} \cap \{1, 4, 5\}$	$\{1\}$	1
(a_1, b_2)	$\{1, 2, 3\} \cap \{2, 3\}$	$\{2, 3\}$	2
(a_2, b_1)	$\{4, 5\} \cap \{1, 4, 5\}$	$\{4, 5\}$	2
(a_2, b_2)	$\{4, 5\} \cap \{2, 3\}$	$\{\}$	0

frag shells



QUERY

A query has the general form
 $\langle a_1, a_2, \dots, a_n : m \rangle$

Each a_i has 3 possible values

Instantiated value (specific; e.g.
 $a_1 = \text{Chicago}$)


Aggregate * function (i.e. don't
care; aggregate over)

Inquire ? function (results
should include all values of this
attribute)

For example:

$\langle 3, ?, ?, *, *, *, I : \text{count} \rangle$

returns a 2-D data cube.



METHOD

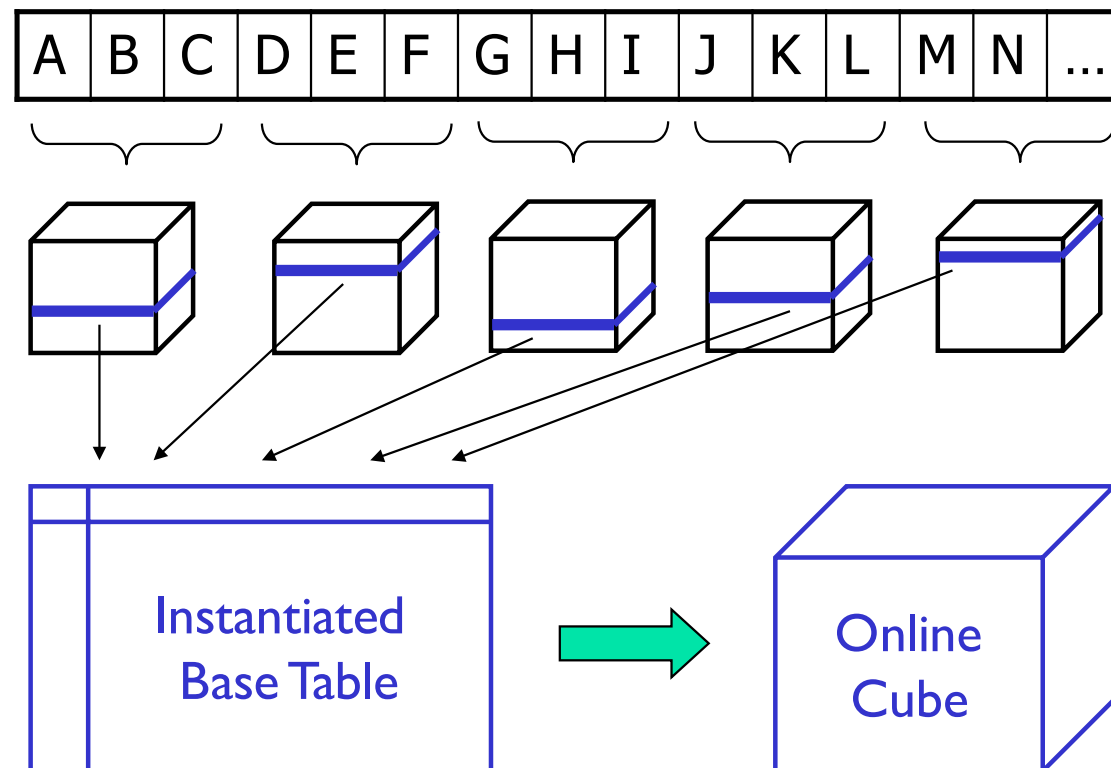

Given the fragment cubes, process a query as follows

Divide the query into fragment, same as the shell


Fetch the corresponding TID list for each fragment from the fragment cube

Intersect the TID lists from each fragment to construct instantiated base table

Compute the data cube using the base table with any cubing algorithm



STORAGE SIZE



Storage grows linearly as the number of dimensions D .

$$O\left(T \left\lceil \frac{D}{F} \right\rceil (2^F - 1)\right)$$

Figure 2: Storage size of shell fragments: (50-C) $T = 10^6$, $C = 50$, $S = 0$, $F = 3$. (100-C) $T = 10^6$, $C = 100$, $S = 2$, $F = 2$.

D denotes the number of dimensions, C is the cardinality of each dimension, T is the number of tuples in the database, F is the size of the shell fragment, I is the number of instantiated dimensions, Q is the number of inquired dimensions, and S is the skew or zipf of the data. **Minimum support level is 1.**

$$O\left(T \left\lceil \frac{D}{F} \right\rceil (2^F - 1)\right)$$




Figure 2: Storage size of shell fragments: (50-C) $T = 10^6$, $C = 50$, $S = 0$, $F = 3$. (100-C) $T = 10^6$, $C = 100$, $S = 2$, $F = 2$.

D denotes the number of dimensions, C is the cardinality of each dimension, T is the number of tuples in the database, F is the size of the shell fragment, I is the number of instantiated dimensions, Q is the number of inquired dimensions, and S is the skew or zipf of the data. **Minimum support level is 1.**

EXPERIMENTS

UCI Forest CoverType data set

54 dimensions, 581K tuples

Shell fragments of size 2 took 33 seconds and 325MB to compute

3-D subquery with 1 instantiated dim: 85ms~1.4 sec.

Longitudinal Study of Vocational Rehab. Data

24 dimensions, 8818 tuples

Shell fragments of size 3 took 0.9 seconds and 60MB to compute

5-D query with 0 instantiated dimensions: 227ms~2.6 sec.




Figure 5: Average computation time per query over 1,000 trials. $\mathcal{T} = 10^6$, $\mathcal{D} = 10$, $\mathcal{C} = 10$, $\mathcal{S} = 0$, $\mathcal{I} = 4$.

what is the value of F , that will lead to smallest fragment cube size?

$$O\left(T\left\lceil\frac{D}{F}\right\rceil(2^F - 1)\right)$$

Why don't we use the smallest value in shell cube design?

$$O\left(T\left\lceil\frac{D}{F}\right\rceil(2^F - 1)\right)$$